

# Use your own data with Azure AI Studio

---

The Azure AI Studio Service enables you to use your own data with the intelligence of the underlying LLM. You can limit the model to only use your data for pertinent topics, or blend it with results from the pre-trained model.

**Duration: 10 to 15 minutes**

## Observe normal chat behavior without adding your own data

Before connecting Azure AI Studio to your data, first observe how the base model responds to queries without any grounding data.

1. Navigate to the **playground**.
2. Enter the following prompts, and observe the output.

I'd like to take a trip to New York. Where should I stay?

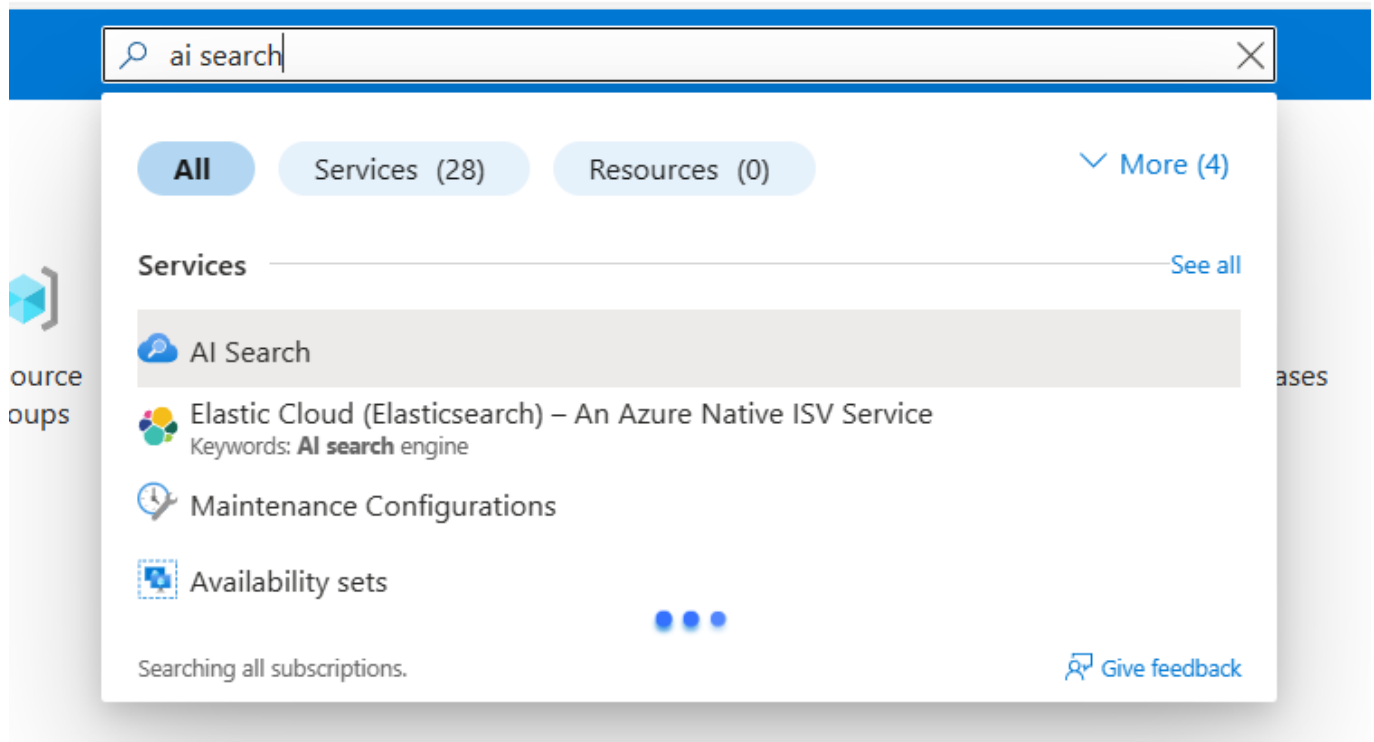
What are some facts about New York?

3. Try similar questions about tourism and places to stay for other locations that will be included in our grounding data, such as London, or San Francisco. You'll likely get complete responses about areas or neighborhoods, and some general facts about the city.

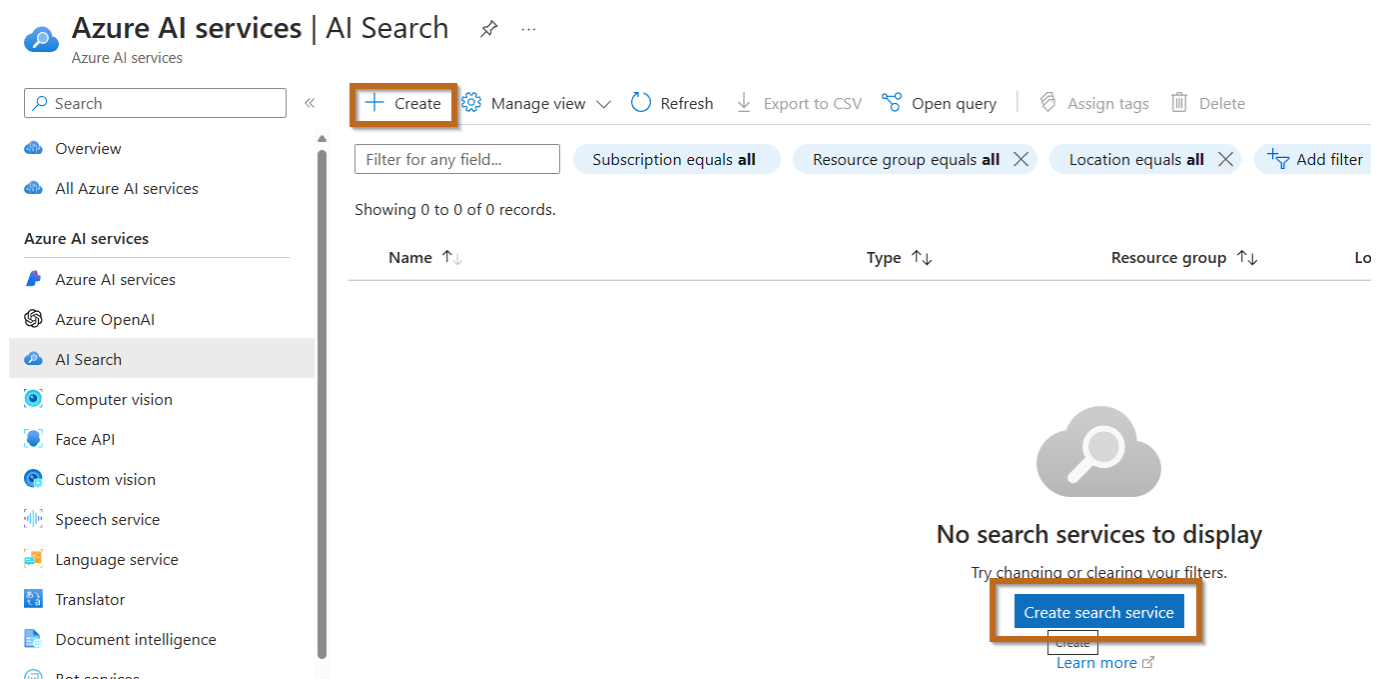
## Create Azure AI Search

Navigate to <https://portal.azure.com/#home>

Search for AI Search.



Create an instance of the service.



**Note:** Azure AI Search is a cloud search service provided by Microsoft Azure that offers AI-powered capabilities to enhance search experiences within applications. It's designed to provide sophisticated search capabilities over large volumes of data, integrating seamlessly with other Azure resources.

Select East US 2 (US Students) or Australia East (Australian Students).

# Create a search service ...

Basics Scale Networking Tags Review + create

## Project details

Subscription \*

Resource Group \*

[Create new](#)

## Instance Details

Service name \* ⓘ

Location \*

Pricing tier \* ⓘ

**Basic**  
2 GB, max 3 replicas, max 1 partitions, max 3 search units  
[Change Pricing Tier](#)

Review + create

Previous

Next: Scale

- **Subscription:** Same subscription as your Azure AI Studio resource
- **Resource group:** Same resource group as your Azure AI Studio resource
- **Service name:** Enter globally unique name
- **Location:** East US 2 (US Students) or Australia East (Australian Students)
- **Pricing tier:** Basic

Review and Create the Service.

# Create a search service ...

✓ Validation Success

Basics   Scale   Networking   Tags   Review + create

## Basics

Subscription	R&D
Resource Group	rg-course-openai-aue-dev-000
Location	Australia East
Service name	(new) srch-davida-dev-aue-001
Pricing tier	basic (2 GB, max 3 replicas, max 1 partitions, max 3 search units)
Estimated cost per month	A\$148.09

## Scale

Replicas	1
Partitions	1

## Networking

Endpoint connectivity (data)	Public
------------------------------	--------

Create

Previous

Next

[Download a template for automation](#)

Once the service is created, navigate to it.

HOME /



## search-service-srch-davida-dev-aue-001 | Overview



Deployment



Delete



Cancel



Redeploy



Download



Overview



Inputs



Outputs



Template



## Your deployment is complete



Deployment name : search-service-srch-davida-dev-a... S

Subscription : R&amp;D C

Resource group : rg-course-openai-aue-dev-000



Deployment details



Next steps


[Go to resource](#)

Give feedback

[Feedback](#)

## Enable Semantic Search.

**Note:** Microsoft is releasing a new method that combines vector, semantic and keyword search.

 **srch-davida-dev-aue-001 | Semantic ranker** ☆ ...  
Search service

- Indexers
- Data sources
- Aliases
- Skillsets
- Debug sessions
- Settings**
  - Semantic ranker**
  - Knowledge Center
  - Keys
  - Scale
  - Search traffic analytics
  - Identity
  - Networking

## Semantic ranker

Semantic ranker uses deep neural networks to provide relevant results an analysis. Additional charges may be applicable.

### Availability

**Free**  
1,000 requests per month  
  
**USD 0.00**/month  
**Select Plan**

**Sta**  
Firs  
\$1.!

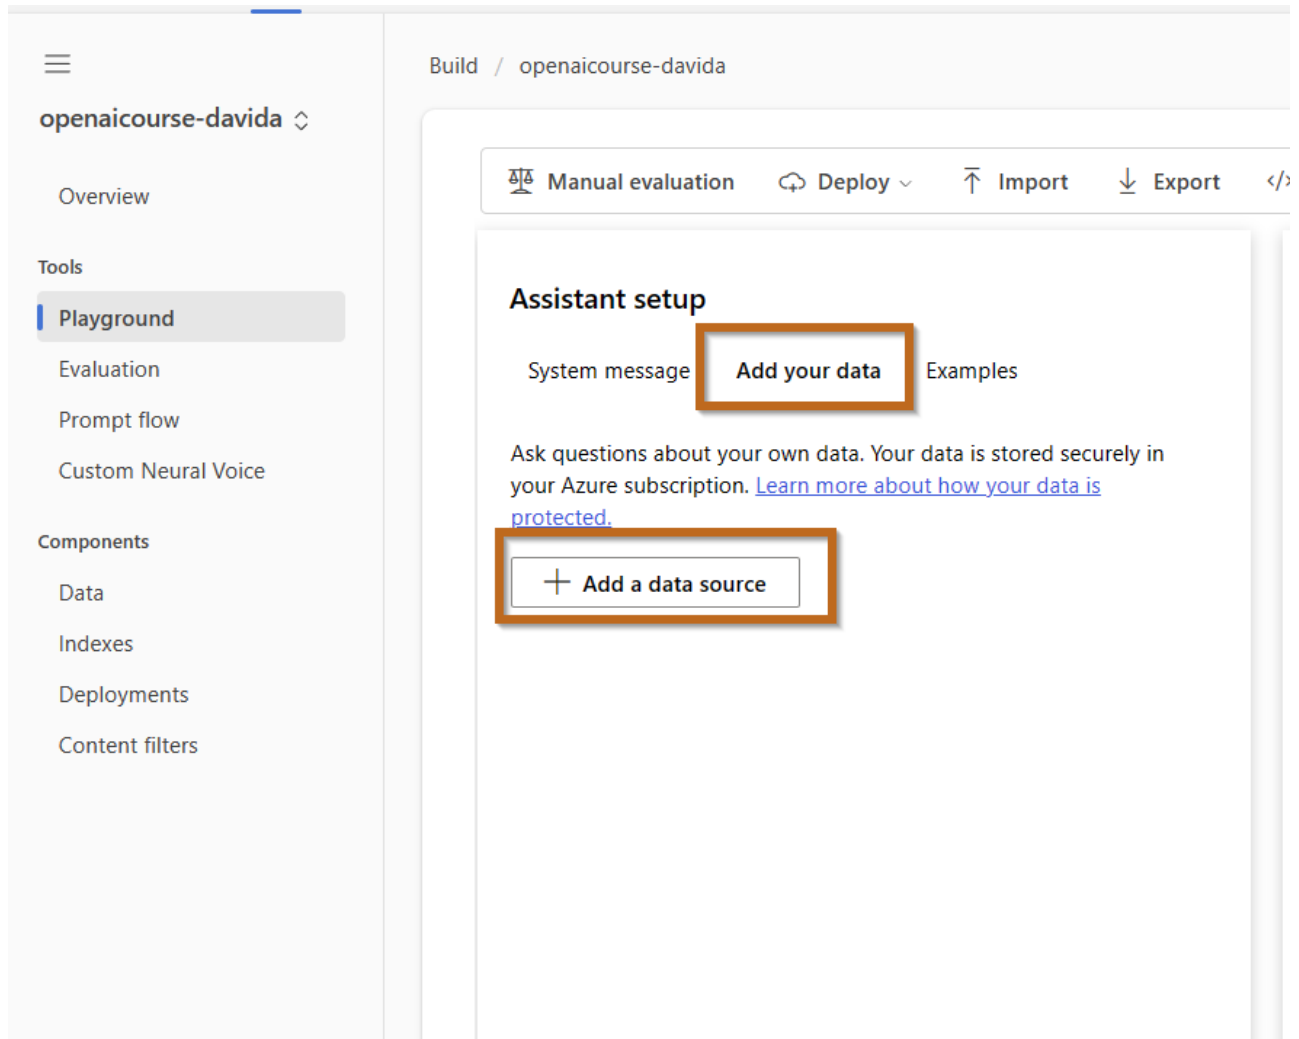
### Learn more

[Availability and pricing](#)

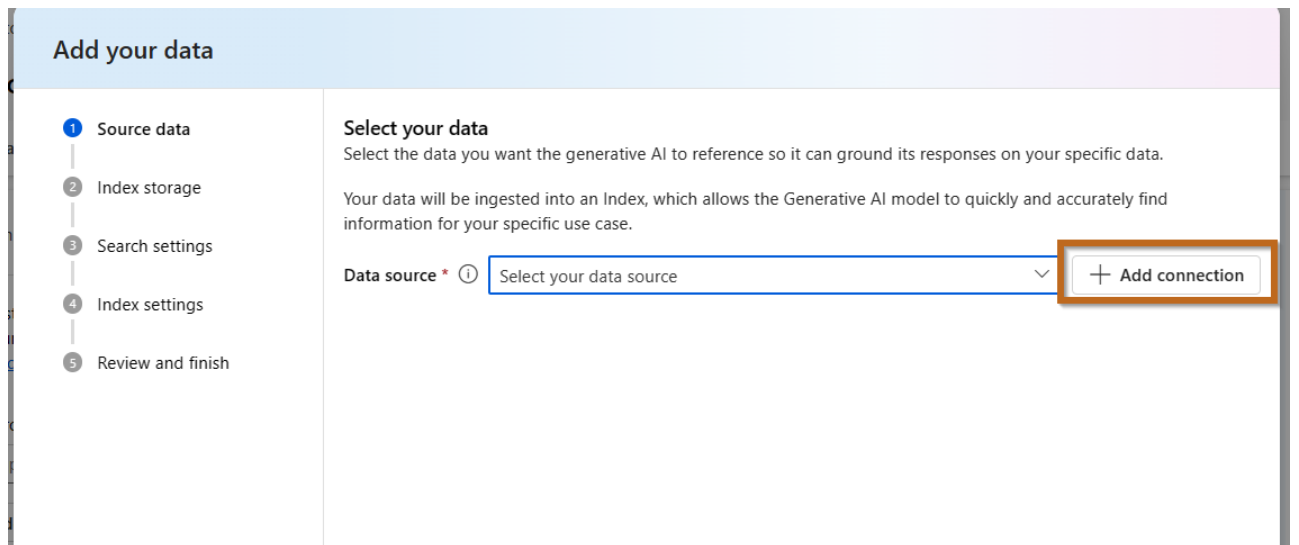
## Connect your data in the chat playground

Next, add your data in the chat playground to see how it responds with your data as grounding

1. Navigate to the **Playground**, and select *Add your data* in the Assistant setup pane.



## 2. Select **Add Connection**.



## 3. Configure the Connection as follow

## Connect an external resource

### Select your service

#### Service \*

Azure Blob Storage



#### Account selection method \*



Manually enter account information

#### Subscription ID \*

David Alzamendi (b28149e7-f6d0-4008-adda-168958c87bc4)



#### Storage account \*

dlsopenaipubdevaue (rg-course-openai-aue-dev-000)



Note: Attaching storage "dlsopenaipubdevaue" can cause higher latencies and bandwidth c...



#### Blob container \*

openaicourse



#### Authentication method \*

Microsoft Entra ID based



#### Connection name \*

storagetravel



#### Access

Project 

Create connection

Cancel

4. Select the previously created AI Search.



### Add your data

- ✓ Source data
- 2 Index storage
- 3 Search settings
- 4 Index settings
- 5 Review and finish

#### Index storage

Select where to store your new index

Index storage \*

Azure AI Search

Select Azure AI Search service \* ⓘ

Connect other Azure AI Search resource

Azure AI Search connection

Subscription \* ⓘ

David Alameedi

Select Azure AI Search service \* ⓘ

srch-davida-dev-aue-001

Previously Created AI Search

Back Next Create Cancel

## 5. Configure the Search Settings as follow

### Add your data

- ✓ Source data
- ✓ Index storage
- 3 Search settings
- 4 Index settings
- 5 Review and finish

#### Configure search settings

Adding vector search supports: Hybrid (vector + keyword search), Hybrid + Semantic (most accurate search results for generative AI applications), Vector, Semantic and Keyword retrieval. Hybrid will be set as default and can be changed at inference time in the playground. Not adding vector search supports: Keyword and Semantic retrieval. Keyword will be set as default and can be changed at inference time in the playground. Adding vector search requires an Azure OpenAI embedding model. [Learn more](#)

Vector settings

☒ Add vector search to this search resource

Azure OpenAI resource \* ⓘ

Default\_AzureOpenAI

☒ I acknowledge that Azure OpenAI embedding model, **text-embedding-ada-002 (Version 2)**, will be deployed if not already.

Back Next Create Cancel

## 6. Configure the Index Settings as Follow

### Add your data

- ✓ Source data
- ✓ Index storage
- ✓ Search settings
- 4 Index settings**
- 5 Review and finish

#### Index settings

Configure your index

Index name \* ⓘ

idx-open-ai-course

Schedule updates \* ⓘ

One time indexing (no scheduled updates) ▾

Virtual machine \* ⓘ

☒ Auto select ☐ Select from recommended options ☐ Select from all options

Back Next Create Cancel

## 7. Review and Finish

### Add your data

- ✓ Source data
- ✓ Index storage
- ✓ Search settings
- ✓ Index settings
- 5 Review and finish**

#### Review and finish

Review the configurations you set for your index

Index name  
idx-open-ai-course

Index storage  
Azure AI Search

Azure AI Search resource  
srch-davida-dev-aue-001

Include vector settings  
Yes

Schedule  
OneTime

Compute  
Serverless compute (Auto select)

Back Create Cancel

8. The information will be captured in a Vector database, you can follow the progress by going to the Index section.

Build / std-000\_ai / std-000-9774 / Indexes

Open project in VS Code (Web)

std-000-9774

Overview

Tools

Playground

Evaluation

Prompt flow

Fine-tuning

Components

Data

Indexes

Deployments

Content filters

AI project settings

Create indexes to customize generative AI responses

+ New index

Refresh

Reset view

Search

Filter

Vector index name	Version	Created by	Data source	Status
idx-open-ai-course		std-000	Azure AI On Your Data	Running

## 9. Click on the Index

### idx-open-ai-course

Status

Running

Refresh

Version

-

Embed with model

No

Source type

Azure AI On Your Data

Vector store

-

Total indexing time

-

Compute

Serverless compute

Created on

Mar 25, 2024, 5:00:39 PM

Created by

std-000

Job details

Test data

Source data

Name	Type	Size
Dubai Brochure.pdf	.pdf	379.51 KE
Las Vegas Brochure.pdf	.pdf	556.56 KE

Status

crack\_chunk\_embed

Created at Mar 25, 2024, 5:00:43 PM

**Note:** When you reach this stage, please inform the instructor. This step will take between 15 and 30 minutes.

## Chat with a model grounded in your data

Now that you've added your data, ask the same questions as you did previously, and see how the response differs.

I'd like to take a trip to New York. Where should I stay?

std-000-9774

Overview

Tools

Playground

Evaluation

Prompt flow

Fine-tuning

Components

Data

Indexes

Deployments

Content filters

AI project settings

Playground

Manual evaluationDeploy to a web appImportExportView CodePrompt samplesCustomize in prompt flowModeChat

System messageAdd your dataExamples

Gain insights into your own data source. Your data is stored securely in your Azure subscription. [Learn more about how your data is protected.](#)

Index:idx-open-ai-course2

Search type:Hybrid (vector + keyword)

[Learn more about different search types](#)

Advanced settings

Remove data source

Replay chatClear chatPlayground SettingsShow JSON

I'd like to take a trip to New York. Where should I stay?

Margie's Travel offers three accommodation options in New York: The Manhattan Hotel, The Grand Central Hotel, and The Park Hotel<sup>1</sup>. The Manhattan Hotel is located in the heart of New York City, within an easy walk to Times Square and Broadway. The Grand Central Hotel is a comfortable mid-town hotel, close to Grand Central Station, the Chrysler Building, and the Empire State Building. The Park Hotel offers luxurious accommodation in upper Manhattan, with views of Central Park<sup>1</sup>. To book your trip to New York, you can visit [www.margiestravel.com](#)<sup>1</sup>.

> 1 references

Type user query here. (Shift + Enter for new line)

DeploymentParameters



Deployment \*gpt-35-turbo

Session settings

Past messages included ⓘ10

What are some facts about New York?

## Chat session

 Clear chat  Playground Settings ☐ Show raw JSON



What are some facts about New York?



New York, often referred to as New York City, is the most populous city in the United States, with an estimated population of 8,550,405 in 2015. The city spans a land area of about 302.6 square miles, making it the most densely populated major city in the United States. Located at the southern tip of the state of New York, it is the center of the New York metropolitan area, one of the most populous urban agglomerations in the world. New York City is a global power city with significant influence on commerce, finance, media, art, fashion, research, technology, education, and entertainment. The city's fast pace has defined the term "New York minute". It is home to the headquarters of the United Nations, making it an important center for international diplomacy. New York has been described as the cultural and financial capital of the world <sup>1</sup>.

1 references >

Type user query here. (Shift + Enter for new line)



You'll notice a very different response this time, with specifics about certain hotels and a mention of Margie's Travel, as well as references to where the information provided came from. If you open the PDF reference listed in the response, you'll see the same hotels as the model provided.

**Note:** **Add your data** is still in preview and might not always behave as expected for this feature, such as giving the incorrect reference for a city not included in the grounding data.

## Navigate to the Data and Index Section

Navigate to the data sections to explore new assets added.

Build / openaia-course-davida / Data

Manage your data

+ New data

Refresh

View options

Search

Name	Version	Created on	Modified on ↓	Type	Created by	Ta
<a href="#">ixcities</a>	1	Nov 19, 2023 2:25 PM	Nov 19, 2023 2:25 PM	Folder	testcourse	
<a href="#">ixdavida</a>	1	Nov 19, 2023 2:25 PM	Nov 19, 2023 2:25 PM	Folder	testcourse	

openaia-course-davida

Overview

Tools

Playground

Evaluation

Prompt flow

Custom Neural Voice

Components

Data

Indexes

Deployments

Content filters

If you click on the data, you will see information about it.

Build / openaia-course-davida / Data / ixcities:1

ixcities:1

Data details

Name

ixcities

Current version

Version: 1 (latest)

+ New version

Latest version

1

Type

Folder

Number of files

1

Total size

636 B

<<

40778850-9021-4658-b200-e95...

azureml

b98697e9-e200-4ba4-870e...

index

MLIndex

openaia-course-davida

Overview

Tools

Playground

Evaluation

Prompt flow

Custom Neural Voice

Components

Data

Indexes

Deployments

Content filters

Navigate to the Index Section.

Confidential - Actionable Data Analytics Pty Ltd (ABN 79 649 102 536) - 14 / 15

☰

openaiaourse-davida

Overview

Tools

Playground

Evaluation

Prompt flow

Custom Neural Voice

Components

Data

Indexes

Deployments

Content filters

Build / openaiaourse-davida / Create indexes to customize generative AI responses

Create indexes to customize generative AI responses

+ New index

↻ Refresh

⚙ View options

🔍 Search

Vector index name	Version	Created by	Data source	Status
ixcities	1	testcourse		✓ Ready
ixdavida	1	testcourse		✓ Ready